

Week 5-8 : Graded Assignment Part 1 - Analysing IoT Data with Spark Sql

Submitted by: Jophy Joseph August 07, 2022

In [1]:

```
!sudo apt update
!apt-get install openjdk-8-jdk-headless -qq > /dev/null
!wget -q https://dlcdn.apache.org/spark/spark-3.3.0/spark-3.3.0-bin-hadoop3.tgz
```

```
Get:1 https://cloud.r-project.org/bin/linux/ubuntu bionic-cran40/ InRelease [3,626 B]
Get:2 http://security.ubuntu.com/ubuntu bionic-security InRelease [88.7 kB]
Hit:3 http://ppa.launchpad.net/c2d4u.team/c2d4u4.0+/ubuntu bionic InRelease
Hit:4 http://archive.ubuntu.com/ubuntu bionic InRelease
Ign:5 https://developer.download.nvidia.com/compute/machine-learning/repos/ubuntu1804/x86_64 InRelease
Get:6 https://developer.download.nvidia.com/compute/cuda/repos/ubuntu1804/x86_64 InRelease [1,581 B]
Hit:7 https://developer.download.nvidia.com/compute/machine-learning/repos/ubuntu1804/x86_64 Release
Get:8 http://archive.ubuntu.com/ubuntu bionic-updates InRelease [88.7 kB]
Get:9 https://developer.download.nvidia.com/compute/cuda/repos/ubuntu1804/x86_64 Packages [902 kB]
Hit:11 http://ppa.launchpad.net/cran/libgit2/ubuntu bionic InRelease
Hit:12 http://ppa.launchpad.net/deadsnakes/ppa/ubuntu bionic InRelease
Get:13 http://security.ubuntu.com/ubuntu bionic-security/main amd64 Packages [2,905 kB]
Get:14 http://archive.ubuntu.com/ubuntu bionic-backports InRelease [74.6 kB]
Hit:15 http://ppa.launchpad.net/graphics-drivers/ppa/ubuntu bionic InRelease
Get:16 http://archive.ubuntu.com/ubuntu bionic-updates/main amd64 Packages [3,336 kB]
Get:17 http://security.ubuntu.com/ubuntu bionic-security/universe amd64 Packages [1,528 kB]
Get:18 http://archive.ubuntu.com/ubuntu bionic-updates/universe amd64 Packages [2,306 kB]
Fetched 11.2 MB in 4s (2,923 kB/s)
Reading package lists... Done
Building dependency tree
Reading state information... Done
25 packages can be upgraded. Run 'apt list --upgradable' to see them.
```

In [2]:

```
!tar xf spark-3.3.0-bin-hadoop3.tgz
!pip install -q findspark
!pip install pyspark
```

```
Looking in indexes: https://pypi.org/simple, https://us-python.pkg.dev/colab-wheels/public/simple/
Collecting pyspark
  Downloading pyspark-3.3.0.tar.gz (281.3 MB)
    |████████████████████████████████████████| 281.3 MB 48 kB/s
Collecting py4j==0.10.9.5
  Downloading py4j-0.10.9.5-py2.py3-none-any.whl (199 kB)
    |████████████████████████████████████████| 199 kB 49.0 MB/s
Building wheels for collected packages: pyspark
  Building wheel for pyspark (setup.py) ... done
  Created wheel for pyspark: filename=pyspark-3.3.0-py2.py3-none-any.whl size=281764026 sha256=1252ffb50a5c44b453f4cdb66a9fba3f5fef5162d0936ba9a307d0e9df73cab7
```

Stored in directory: /root/.cache/pip/wheels/7a/8e/1b/f73a52650d2e5f337708d9f6a1750d451a7349a867f928b885
Successfully built pyspark
Installing collected packages: py4j, pyspark
Successfully installed py4j-0.10.9.5 pyspark-3.3.0

```
In [3]: import os
os.environ["JAVA_HOME"] = "/usr/lib/jvm/java-8-openjdk-amd64"
os.environ["SPARK_HOME"] = "/content/spark-3.3.0-bin-hadoop3"
```

```
In [4]: import findspark
findspark.init()
findspark.find()
```

Out[4]: '/content/spark-3.3.0-bin-hadoop3'

```
In [5]: from pyspark.sql import DataFrame, SparkSession
from typing import List
import pyspark.sql.types as T
import pyspark.sql.functions as F

spark = SparkSession \
    .builder \
    .appName("Part-1: Working with SparkSQL") \
    .getOrCreate()

spark
```

Out[5]: **SparkSession - in-memory**

SparkContext

[Spark UI](#)

Version	v3.3.0
Master	local[*]
AppName	Part-1: Working with SparkSQL

Task 1: Read the data into a Dataframe.

```
In [6]: from google.colab import drive
drive.mount('/content/drive')
```

Mounted at /content/drive

```
In [7]: file_location = "/content/drive/My Drive/data/iot_devices.json"
file_type = "json"

# CSV options
infer_schema = True
first_row_is_header = True
delimiter = ","

# The applied options are for CSV files. For other file types, these will be ignored.
iot_df = spark.read.format(file_type) \
    .option("inferSchema", infer_schema) \
    .option("header", first_row_is_header) \
    .option("sep", delimiter) \
    .load(file_location)
```

```
In [8]: iot_df.show(20, False)
```

	battery_level	c02_level	cca2	cca3	cn	device_id	device_name	humidity	ip	latitude	lcd	longitude	scale	temp	timestamp
8	868		US	USA	United States	1	meter-gauge-1xbYRYcj	51	68.161.225.1	38.0	green	-97.0	Celsius	34	1458444054093
7	1473		NO	NOR	Norway	2	sensor-pad-2n2Pea	70	213.161.254.1	62.47	red	6.15	Celsius	11	1458444054119
2	1556		IT	ITA	Italy	3	device-mac-36TWSKiT	44	88.36.5.1	42.83	red	12.83	Celsius	19	1458444054120
6	1080		US	USA	United States	4	sensor-pad-4mzWkz	32	66.39.173.154	44.06	yellow	-121.32	Celsius	28	1458444054121
4	931		PH	PHL	Philippines	5	therm-stick-5gimpUrBB	62	203.82.41.9	14.58	green	120.97	Celsius	25	1458444054122
3	1210		US	USA	United States	6	sensor-pad-6a17RTAobR	51	204.116.105.67	35.93	yellow	-85.46	Celsius	27	1458444054122
3	1129		CN	CHN	China	7	meter-gauge-7GeDoanM	26	220.173.179.1	22.82	yellow	108.32	Celsius	18	1458444054123
0	1536		JP	JPN	Japan	8	sensor-pad-8xUD6pzsQI	35	210.173.177.1	35.69	red	139.69	Celsius	27	1458444054123
3	807		JP	JPN	Japan	9	device-mac-9GcjZ2pw	85	118.23.68.227	35.69	green	139.69	Celsius	13	1458444054124

7	1470	US	USA	United States	10	sensor-pad-10BsywSYUF	56	208.109.163.218	33.61	red	-111.89	Celsius	26	1458444054125
3	1544	IT	ITA	Italy	11	meter-gauge-11dlMTZty	85	88.213.191.34	42.83	red	12.83	Celsius	16	1458444054125
0	1260	US	USA	United States	12	sensor-pad-12Y2kIm0o	92	68.28.91.22	38.0	yellow	-97.0	Celsius	12	1458444054126
6	1007	IN	IND	India	13	meter-gauge-13GrojanSGBz	92	59.144.114.250	28.6	yellow	77.2	Celsius	13	1458444054127
1	1346	NO	NOR	Norway	14	sensor-pad-14QL93sBR0j	90	193.156.90.200	59.95	yellow	10.75	Celsius	16	1458444054127
9	1259	US	USA	United States	15	device-mac-15se6mZ	70	67.185.72.1	47.41	yellow	-122.0	Celsius	13	1458444054128
4	1425	US	USA	United States	16	sensor-pad-16aXmIJZtd0	53	68.85.85.106	38.0	red	-97.0	Celsius	15	1458444054128
0	1466	US	USA	United States	17	meter-gauge-17zb8Fghh1	98	161.188.212.254	39.95	red	-75.16	Celsius	31	1458444054129
4	1096	CN	CHN	China	18	sensor-pad-18XULN9Xv	25	221.3.128.242	25.04	yellow	102.72	Celsius	31	1458444054130
9	1531	US	USA	United States	19	meter-gauge-19eg1BpfCO	75	64.124.180.215	38.0	red	-97.0	Celsius	29	1458444054130
7	1155	US	USA	United States	20	sensor-pad-20gFNfBgqr	33	66.153.162.66	33.94	yellow	-78.92	Celsius	10	1458444054131

+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+
+
only showing top 20 rows

Task 2. Convert the Dataframe into a temporary view called `iot`.

In [9]:

```
iot_df.createOrReplaceTempView('iot')
```

In [10]:

```
spark.sql("select * from iot").show(5)
```

battery_level	c02_level	cca2	cca3	cn	device_id	device_name	humidity	ip	latitude	lcd	longitude	scale	temp	timestamp
8	868	US	USA	United States	1	meter-gauge-1xbYRYcj	51	68.161.225.1	38.0	green	-97.0	Celsius	34	1458444054093
7	1473	NO	NOR	Norway	2	sensor-pad-2n2Pea	70	213.161.254.1	62.47	red	6.15	Celsius	11	1458444054119
2	1556	IT	ITA	Italy	3	device-mac-36TWSKiT	44	88.36.5.1	42.83	red	12.83	Celsius	19	1458444054120
6	1080	US	USA	United States	4	sensor-pad-4mzWkz	32	66.39.173.154	44.06	yellow	-121.32	Celsius	28	1458444054121
4	931	PH	PHL	Philippines	5	therm-stick-5gimp...	62	203.82.41.9	14.58	green	120.97	Celsius	25	1458444054122

only showing top 5 rows

```
In [11]: iot_df.printSchema()
```

```
root
|-- battery_level: long (nullable = true)
|-- c02_level: long (nullable = true)
|-- cca2: string (nullable = true)
|-- cca3: string (nullable = true)
|-- cn: string (nullable = true)
|-- device_id: long (nullable = true)
|-- device_name: string (nullable = true)
|-- humidity: long (nullable = true)
|-- ip: string (nullable = true)
|-- latitude: double (nullable = true)
|-- lcd: string (nullable = true)
|-- longitude: double (nullable = true)
|-- scale: string (nullable = true)
|-- temp: long (nullable = true)
|-- timestamp: long (nullable = true)
```

Task 3. Count how many devices are there from each country and display the output.

```
In [12]: spark.sql("select cca3 as Country_name, count(distinct(device_id)) as Number_of_devices from iot group by cca3 order by Number_of_devices desc").show()
```

Country_name	Number_of_devices
USA	70405
CHN	14455
JPN	12100
KOR	11879
DEU	7942
GBR	6486
CAN	6041
RUS	5989
FRA	5305
BRA	3224
AUS	3119
ITA	2915
SWE	2880
POL	2744
NLD	2488
ESP	2310
TWN	2128
IND	1867
CZE	1507
NOR	1487

only showing top 20 rows

Task 4. Display all the countries whose carbon dioxide level is more than 1400. Sort the output in descending order.

```
In [13]: spark.sql("select Cca2 as Country_Code, Cca3 as Country_name, sum(c02_level) as CO2_level  from iot  where c02_level > 1400  group by Cca2, Cca3 order by CO2_lev
```

Country_Code	Country_name	CO2_level
US	USA	26242891
CN	CHN	5424312
KR	KOR	4415118
JP	JPN	4399107
DE	DEU	2950796
GB	GBR	2488574
CA	CAN	2343270
RU	RUS	2262936
FR	FRA	2030583
BR	BRA	1284892
AU	AUS	1153899
SE	SWE	1086146
IT	ITA	1070505
PL	POL	995721
NL	NLD	970297
ES	ESP	878143
TW	TWN	813253
IN	IND	666101
NO	NOR	598248
UA	UKR	559605

only showing top 20 rows

Task 5. Select all countries' devices with high-levels of C02 and group by cca3 and order by device_ids (Hint: For high CO2 level, the LCD status will be RED).

```
In [14]: spark.sql("select cca3,cn,device_id from iot where lcd = 'red' group by cca3,cn,device_id order by device_id ").show()
```

cca3	cn	device_id
NOR	Norway	2
ITA	Italy	3
JPN	Japan	8
USA	United States	10
ITA	Italy	11
USA	United States	16

USA	United States	17
USA	United States	19
JPN	Japan	22
CAN	Canada	24
KOR	Republic of Korea	27
KOR	Republic of Korea	28
UKR	Ukraine	47
SWE	Sweden	53
USA	United States	54
USA	United States	57
USA	United States	64
CZE	Czech Republic	66
IND	India	77
KOR	Republic of Korea	78

+-----+

only showing top 20 rows

Task 6. Find out all devices in countries whose batteries need replacements.

In [15]:

```
# Assumption - devices with battery level < 3 would need replacement
spark.sql("select cca3, device_id, battery_level from iot where battery_level < 3  order by battery_level ").show()
```

cca3	device_id	battery_level
JPN	106121	0
DEU	106203	0
AUS	106061	0
RUS	106125	0
AUS	106067	0
VNM	106139	0
USA	106076	0
CAN	106152	0
USA	106088	0
FRA	106162	0
DEU	106105	0
USA	106168	0
GBR	106112	0
USA	106172	0
CHN	106199	0
KOR	106177	0
MYS	106085	0
ROU	106182	0
USA	106110	0
LBY	106195	0

+-----+

only showing top 20 rows